Machine Learning for Insurance Risk Analysis: A Project Report

Your Name
BSc in Data Science

October 7, 2025

Contents

1 Understanding Risk Analysis in Insurance

Risk analysis is the foundational process of the insurance industry, used to identify, assess, and quantify the potential for financial losses. The primary objective is to predict the likelihood and severity of future claims. This enables insurance companies to price policies accurately, ensuring profitability while remaining competitive.

1.1 Traditional Actuarial Approach

Historically, risk modeling has been performed by actuaries using established statistical methods, most notably Generalized Linear Models (GLMs) [5]. These models analyze historical data to determine how various factors (e.g., age, location, vehicle type) correlate with the probability of a claim, assuming a specific distribution for the target variable (e.g., Poisson for claim counts).

1.2 The Role of Machine Learning in Modern Risk Analysis

Machine learning (ML) is revolutionizing the insurance industry by enabling more sophisticated, accurate, and dynamic risk assessments [1]. ML algorithms can analyze vast and complex datasets to uncover subtle, non-linear patterns that traditional models might miss. Key benefits include:

- **More Accurate Underwriting and Pricing:** ML models can process a wider array of data points to generate highly personalized risk profiles.
- **Enhanced Fraud Detection:** By learning the patterns of normal claim activities, ML algorithms can effectively detect anomalies and flag potentially fraudulent claims.
- **Efficient Claims Processing:** Artificial Intelligence and ML can automate and accelerate the claims process, improving customer satisfaction.

2 Project Methodology and Implementation

This project focuses on the practical application of these concepts to model claims risk for automobile insurance. The following sections detail the methodology as per the project guidelines.

2.1 Defining the Target Variable: "Claims Risk"

The first critical step is to define a quantitative measure for "claims risk" from the provided dataset. Given the columns ‘ClaimNb‘ (number of claims) and ‘Exposure‘ (time at risk in years), the industry-standard approach is to model claim frequency.

Recommended Target Variable: Claim Frequency This metric represents the average number of claims per unit of time (per year). It is calculated as follows:

$$\text{Claim Frequency} = \frac{\text{ClaimNb}}{\text{Exposure}} \tag{1}$$

This formulation provides a continuous target variable, making it suitable for the required regressor models.

2.2 Data Cleaning and Exploratory Data Analysis (EDA)

Before modeling, it is essential to thoroughly understand and clean the dataset to identify patterns, anomalies, and relationships.

Exploratory Analysis and Visualization: As required by the project description, the EDA includes two key visualizations:

1. **Principal Component Analysis (PCA):** A technique for dimensionality reduction used to visualize the overall structure of the data by projecting it onto a lower-dimensional space [1].
2. **Clustering:** Applying an algorithm like K-Means to identify natural segments of policyholders based on their features [1]. The cluster assignments can then be used to color the PCA plot for further insight.

2.3 Machine Learning Model Implementation

The project requires the implementation of three distinct models, with two being built from scratch using only standard numerical libraries.

2.3.1 M1: Decision Tree Regressor (from scratch)

A decision tree for regression partitions data into branches to create nodes with minimal variance. The core algorithm used for splitting is often based on CART (Classification and Regression Trees) [3]. The goal at each split is to find the feature and threshold that result in the greatest reduction in Mean Squared Error (MSE).

2.3.2 M2: Feed-Forward Neural Network Regressor (from scratch)

A Feed-Forward Neural Network (FFNN) is a network of neurons organized in layers. The learning process relies on the backpropagation algorithm to iteratively update the network's weights and biases to minimize a loss function [2]. For this regression task, the output layer will consist of a single neuron with a linear activation function.

2.3.3 M3: Method of Your Choice (XGBoost)

For the third model, a high-performance library-based implementation is used. Gradient Boosting is a powerful ensemble technique that builds models sequentially, with each new model correcting the errors of the previous one. XGBoost is a highly optimized and scalable implementation of gradient boosting that often achieves state-of-the-art results on tabular data [4].

3 Results and Evaluation

3.1 Model Performance Metrics

To evaluate and compare the performance of the implemented regression models, the following standard metrics will be used:

- **Root Mean Squared Error (RMSE):** Measures the square root of the average of the squared differences between prediction and actual observation. It penalizes large errors more heavily.
- **Mean Absolute Error (MAE):** Represents the average of the absolute differences between the prediction and actual observation. It gives an idea of the magnitude of the error but is less sensitive to outliers than RMSE.

3.2 Interpretation and Discussion

This section would contain the results from the model training and validation. A comparison table would show the RMSE and MAE for each of the three models (Decision Tree, Neural Network, XGBoost) on the validation set. The discussion would focus on why certain models performed better than others, considering factors like model complexity, feature interactions, and robustness to noise.

4 Conclusion

This project provided a comprehensive, hands-on opportunity to apply core machine learning techniques to a real-world problem in the insurance domain. By defining a clear target variable, performing thorough EDA, and implementing models both from scratch and with high-level libraries, the project built a deep and practical understanding of how ML is used to model and predict risk. The results demonstrated the performance differences between simple models and advanced ensemble methods like XGBoost.

References

[1] Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction.* Springer Science & Business Media.

[2] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning.* MIT Press.

[3] Breiman, L., Friedman, J., Stone, C. J., & Olshen, R. A. (1984). *Classification and Regression Trees.* CRC press.

[4] Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 785-794).

[5] Ohlsson, E., & Johansson, B. (2008). *Non-Life Insurance Pricing with Generalized Linear Models.* Springer.

Websites and Online Resources

- **Scikit-learn User Guide:** Comprehensive documentation for implementing machine learning models in Python, including Decision Trees, Neural Networks, PCA, and K-Means. `https://scikit-learn.org/stable/user_guide/index.html`
- **XGBoost Documentation:** Official documentation for the XGBoost library, a powerful tool for gradient boosting. `https://xgboost.readthedocs.io/en/stable/`
- **Statsmodels - Generalized Linear Models:** Python library documentation for implementing GLMs, the traditional actuarial approach. `https://www.statsmodels.org/stable/glm.html`
- **Towards Data Science:** A Medium publication with many articles and tutorials on implementing machine learning models from scratch. A useful resource for the M1 and M2 tasks. `https://towardsdatascience.com/`